

Docker installieren

Install Docker Desktop: <https://hub.docker.com/editions/community/docker-ce-desktop-windows/>

Login Docker

Switch to Linux Containers

Images holen

<code>docker login</code>
<code>docker pull tinkertop/gremlin-console</code>
<code>docker pull mongo:latest</code>
<code>docker pull node:latest</code>
<code>docker pull bitnami/cassandra:latest</code>

Azure Ressourcen installieren

Install latest CLI

<https://docs.microsoft.com/en-us/cli/azure/install-azure-cli-windows?view=azure-cli-latest>

<code>Invoke-WebRequest -Uri https://aka.ms/installazurecliwindows -OutFile .\AzureCLI.msi; Start-Process msixexec.exe -Wait -ArgumentList '/I AzureCLI.msi /quiet'; rm .\AzureCLI.msi</code>

Resource Group: GAB2020

Create CosmosDB Accounts:

- gab2020sql (CORESQL) Notebook Preview ON, NO FREE TIER,

<code>az login</code>	
<code>az account set -s <yourSubscriptionID></code>	
Create Accounts	
<code>az cosmosdb create --resource-group GAB2020 --name gabus2020sql --locations regionName=westus</code>	Core SQL

az cosmosdb create --resource-group GAB2020 --name gabus2020cas --locations regionName=westus --capabilities EnableCassandra	Cassandra
az cosmosdb create --resource-group GAB2020 --name gabus2020graph --locations regionName=westus --capabilities EnableGremlin	Gremlin
az cosmosdb create --resource-group GAB2020 --name gabus2020mongo --locations regionName=westus --kind MongoDB	Mongo DB - Wire Protokoll 3.2
Muss im Portal angelegt werden (keine CLI/ARM Unterstützung zur Zeit) - gabus2020mongo	Mongo DB - Wire Protokoll 3.6

Git Repository holen

C:

cd \Source\2020\Events\

git clone <https://github.com/apollak/GAB2020-CosmosDB-AT.git> GAB2020

Endpoint URL und Passwort für Core SQL setzen

// SETX setzt das Maschinenweit / Erfordert neue Instanz von Terminal! (nicht nur TAB!)

```
setx EndpointUrl "https://gabus2020sql.documents.azure.com:443/"
```

```
setx PrimaryKey "<your primary key>"
```

Core SQL – Demos

Basic Demo

Terminal

```
cd C:\Source\2020\Events\GAB2020\CoreSQL\Basic
```

```
setx EndpointUrl "<Your_Azure_Cosmos_account_URI>"
```

```
setx PrimaryKey "<Your_Azure_Cosmos_account_PRIMARY_KEY>"
```

```
code .
```

Spatial Demo

Terminal

```
cd C:\Source\2020\Events\GAB2020\CoreSQL\Spatial

setx EndpointUrl "<Your_Azure_Cosmos_account_URI>"
setx PrimaryKey "<Your_Azure_Cosmos_account_PRIMARY_KEY>"

code .
```

Graph-Demo

<https://portal.azure.com>

Select: **gabus2020graph**

DataExplorer

- Create Database 400 RU/s (**graphdb**)
- Create Collection (**thehobbit**) - Partition Key (/p)

```
cd C:\Source\2020\Events\GAB2020\Gremlin\msgraph

setx GraphEndpointUrl "<your graph db account>.gremlin.cosmos.azure.com"
setx GraphPrimaryKey "<your primary key>"
```

Terminal

```
cd C:\Source\2020\Events\GAB2020\Gremlin
mkdir conf
code remote.yaml
```

remote.yaml (<https://docs.microsoft.com/en-us/azure/cosmos-db/create-graph-gremlin-console>)

```
hosts: [your_database_server.gremlin.cosmos.azure.com]
port: 443
username: /dbs/your_database_account/colls/your_collection
password: your_primary_key
connectionPool: {
  enableSsl: true
}
serializer: { className:
org.apache.tinkerpop.gremlin.driver.ser.GraphSONMessageSerializerV2d0, config: {
serializeResultToString: true }}
```

Open Terminal

C:\Source\2020\Events\GAB2020\Gremlin	
docker run -it --volume \${PWD}:/pc --name graph tinkerpops/gremlin-console bash	Beim ersten Aufruf
docker start -i graph	Danach

Alternativ: Cd C:\Source\2020\Events\GAB2020\Gremlin\msGraph => Simple REPL verwenden

Connect to Graph
:remote connect tinkerpops.server /pc/conf/remote.yaml
:remote console
g.addV('person').property('name', 'Bilbo Beutlin').property('p', 'ring');
g.addV('person').property('name', 'Gandalf').property('p', 'ring'); g.addV('person').property('name', 'Thorin').property('p', 'ring'); g.addV('person').property('name', 'Dwalin').property('p', 'ring'); g.addV('person').property('name', 'Balin').property('p', 'ring'); g.addV('person').property('name', 'Kili').property('p', 'ring'); g.addV('person').property('name', 'Fili').property('p', 'ring'); g.addV('person').property('name', 'Dori').property('p', 'ring'); g.addV('person').property('name', 'Nori').property('p', 'ring'); g.addV('person').property('name', 'Ori').property('p', 'ring'); g.addV('person').property('name', 'Oin').property('p', 'ring'); g.addV('person').property('name', 'Gloin').property('p', 'ring'); g.addV('person').property('name', 'Bifur').property('p', 'ring'); g.addV('person').property('name', 'Bofur').property('p', 'ring'); g.addV('person').property('name', 'Bombur').property('p', 'ring');
g.V()
Bilbo connecting (pre corona)
g.V().has('name', 'Bilbo Beutlin').addE('knows').to(g.V().has('name', 'Gandalf'))
g.V().has('name', 'Bilbo Beutlin').addE('knows').to(g.V().has('name', 'Thorin'))
Gandalf connects to everybody
g.V().has('name', 'Gandalf').as('gandalf').V().hasLabel('person').where(neq('gandalf')).as('p').V().has('name', 'Gandalf').addE('knows').to(select('p'))
Add Places
g.addV('place').property('name', 'Hobbithöhle').property('p', 'ring'); g.addV('place').property('name', 'Gasthaus Zum grünen Drachen').property('p', 'ring'); g.addV('place').property('name', 'Troll Lagerplatz').property('p', 'ring'); g.addV('place').property('name', 'Trollhöhle').property('p', 'ring'); g.addV('place').property('name', 'Rivendell').property('p', 'ring');

<pre>g.addV('place').property('name','Elronds Haus').property('p','ring'); g.addV('place').property('name','Orkhöhle').property('p','ring');</pre>
Add Paths
<pre>g.V().has('name','Hobbithöhle').addE('path').to(g.V().has('name','Gasthaus Zum grünen Drachen')).property('weight',2.0); g.V().has('name','Gasthaus Zum grünen Drachen').addE('path').to(g.V().has('name','Troll Lagerplatz')).property('weight',4.0); g.V().has('name','Troll Lagerplatz').addE('path').to(g.V().has('name','Trollhöhle')).property('weight',1.0); g.V().has('name','Troll Lagerplatz').addE('path').to(g.V().has('name','Rivendell')).property('weight',3.0); g.V().has('name','Rivendell').addE('path').to(g.V().has('name','Elronds Haus')).property('weight',1.0); g.V().has('name','Elronds Haus').addE('path').to(g.V().has('name','Orkhöhle')).property('weight',5.0); g.V().has('name','Gasthaus Zum grünen Drachen').addE('path').to(g.V().has('name','Rivendell')).property('weight',4.0); g.V().has('name','Gasthaus Zum grünen Drachen').addE('path').to(g.V().has('name','Orkhöhle')).property('weight',8.0);</pre>
<pre>g.V(); g.E();</pre>
<pre>g.V().has('name','Hobbithöhle').repeat(outE().inV().simplePath()).until(has('name','Orkhöhle')).pa th().by(coalesce(values('weight'),constant(0.0))).map(unfold().sum());</pre>
<pre>g.V().has('name','Hobbithöhle').repeat(outE().inV().simplePath()).until(has('name','Orkhöhle')).pa th().by(coalesce(values('name'),constant(0.0)));</pre>
<pre>g.V().has('name','Hobbithöhle').repeat(outE().inV().simplePath()).until(has('name','Orkhöhle')).pa th().by(coalesce(values('name','weight'),constant(0.0)));</pre>
:remote console
:quit

MongoDB

<https://portal.azure.com>

Select: **gabus2020mongo**

Select: **ConnectionString => Primary Key**

Adapt primary-template.key file and copy to primary.key

Terminal

cd C:\Source\2020\Events\GAB2020\Mongo	
docker run -it --volume \${PWD}:/home --name mongo mongo bash	Beim ersten Aufruf
docker start -i mongo	Danach

Connect to MongoDB
cd /home
./primary.key
mongo gabus2020mongo.mongo.cosmos.azure.com:10255 -u gabus2020mongo -p \$primarykey -tls --tlsAllowInvalidCertificates

News

- Azure Cosmos DB's API for MongoDB unterstützt **Server Version 3.6**
 - Verbesserte Latency bei Group, Count and Skip-Limit
 - Compound indexes
 - ChangeFeed support via ChangeStream API
 - Creating unsharded collections under databases with throughput
 - Aggregation pipeline stages/operators

Unsharded Collections (Database Throughput does no longer require a partition key)

use mydb
Create explicit Throughput
db.runCommand({customAction: "CreateDatabase", offerThroughput: 400});
Create explicit Throughput on collection (based on DB-Value!)
db.createCollection("democol");
Create implicit Throughput on collection (uses db Rus)
db.runCommand({customAction: "CreateCollection", collection: "democol2"});
Create explicit Throughput on collection based on parameter!
db.runCommand({customAction: "CreateCollection", collection: "democol3", offerThroughput: 1000 });
db.democol.insert({a:1});
Creates implicit throughput on collection (uses db Rus)
db.runCommand({customAction: "CreateCollection", collection: "democol1", shardKey: "a.b" });
db.democol1.insert({a:{ b: 'Hugo' }});
db.democol1.insert({a:{ c: 'Hugo' }});
Works in previous collection because there is no shard key

```
db.democol.insert({a:{ c: 'Hugo' }});
```

Compound Indexes (democol2)

Compound indexes give a huge performance boost for queries with:

- **Multiple properties in the Sort()** definition
- Find() and Sort() where the **property in the Find() definition is part of Sort()**
- Find() with multiple properties filter where at least one property filter is equality.(e.g. db.test.find({a:1, b: {\$gte: 5}, b: {\$lte: 7}});)

```
db.democol2.insertMany( [  
  { a:1, b: 5, i:"Hansi" },  
  { a:1, b: 2, i:"Susi" },  
  { a:2, b: 10, i:"Maria" },  
  { a:1, b: 7, i:"Karl" },  
  { a:1, b: 9, i:"Cornelia" },  
]);  
  
db.democol2.createIndex({a:1, b:1})  
db.democol2.find({a:1, b: {$gte: 5}, b: {$lte: 7}})  
db.democol2.find({a:1, b: {$gte:5, $lte: 7}})
```

ChangeFeed via ChangeStream API (democol)

<https://docs.mongodb.com/manual/reference/method/db.collection.watch/#db.collection.watch>

Watch Changefeed for documents that have property a set!

```
var watchCursor = db.democol.watch(  
  [  
    {  
      $match: {  
        $and: [  
          { "fullDocument.a": 1 },  
          { "operationType": { $in: ["insert", "update", "replace"] } }  
        ]  
      }  
    },  
    { $project: { "_id": 1, "fullDocument": 1, "ns": 1, "documentKey": 1 } }  
  ],  
  { fullDocument: "updateLookup" });
```

watchCursor.hasNext()
db.democol.insert({a:1,b:"Hello Change"})
watchCursor.hasNext()
watchCursor.next()
watchCursor.close()
while (!watchCursor.isExhausted()){ if (watchCursor.hasNext()){ printjson(watchCursor.next()); } }

Faster Aggregation pipeline stages/operators (democol2)

use mydb
Example-Data
db.democol2.insertMany([{ animal:"hedgehog", color:"brown", a:1 }, { animal:"hedgehog", color:"indigo", a:2 }, { animal:"hedgehog", color:"blue", a:3 }, { animal:"sheep", color:"blue", a:5 }, { animal:"sheep", color:"yellow", a:6 }, { animal:"sheep", color:"brown", a:7 }, { animal:"sheep", color:"orange", a:8 }, { animal:"ape", color:"orange", a:4 }, { animal:"ape", color:"black", a:3 }]);
Sample Aggregates
db.democol2.aggregate({\$sortByCount:"\$animal"})
db.democol2.aggregate([{\$group: {_id:"\$color", count:{\$sum:1},a:{\$min:"\$a"} }},{\$sort:{count:-1}}])
Delete items from collection
db.democol2.deleteMany({});
exit

TTL (Time To Live Support) (democol)

Default 10 Seconds
db.democol.createIndex({"_ts":1}, {expireAfterSeconds: 10})
db.democol.insertMany([{ name:"TTL Hansi" },

<code>{ name:"TTL Susi", ttl:6 }</code>
<code>]);</code>
<code>db.democol.find({});</code>
Drop Index
<code>db.democol.dropIndex({"_ts":1})</code>
No Default
<code>db.democol.createIndex({"_ts":1}, {expireAfterSeconds: -1})</code>
<code>db.democol.insertMany([</code> <code> { name:"TTL Hansi" },</code> <code> { name:"TTL Susi", ttl:6 }</code> <code>]);</code>
<code>db.democol.find({});</code>
<code>db.democol.find({});</code>

Cassandra - Demo

Start Node Environment

2 Terminal Sessions öffnen und in beiden in das nachfolgende Verzeichnis wechseln	
<code>cd C:\Source\2020\Events\GAB2020\Cassandra</code>	
Terminal-1 (Node Environment)	
<code>docker run -it -p 3000:3000 --volume \${PWD}:/home --name nodedev node bash</code>	#Initial run
<code>docker start -i nodedev</code>	#restart run
<code>docker container rm nodedev</code>	#remove container

Cassandra Simple Demo

Terminal-1 Node Environment
<code>cd /home</code>
<code>git clone https://github.com/doanduyhai/Cassandra-NodeJS-Demo.git</code>
<code>git clone https://github.com/Azure-Samples/azure-cosmos-db-cassandra-nodejs-getting-started.git</code>
<code>cp config.js Cassandra-NodeJS-Demo/config.js</code>
<code>cp config.js azure-cosmos-db-cassandra-nodejs-getting-started/config.js</code>
<code>cd Cassandra-NodeJS-Demo</code>
<code>npm install</code>
<code>openssl s_client -connect gabus2020cas.cassandra.cosmos.azure.com:10350 openssl x509 -out gabus2020cas.pem</code>

Terminal-2 Node Environment
code .
Show config-Template.js
Navigate to /Cassandra-NodeJS-Demo/src/Ingestionjs

Navigate to: <https://portal.azure.com>

Open **gabus2020cas**

Click DataExplorer

New Keyspace: **nodejs_demo**

Click **Quickstart - NodeJS**

Change Code: src/Ingestion.js and src/Reading.js
Replace the following line of code: <pre>var client = new cassandra.Client({contactPoints: ['localhost']});</pre>
WITHOUT SSL-VALIDATION
<pre>var config = require('../config'); process.env.NODE_TLS_REJECT_UNAUTHORIZED = "0"; const authProviderLocalCassandra = new cassandra.auth.PlainTextAuthProvider(config.username, config.password); const client = new cassandra.Client({contactPoints: [config.contactPoint], authProvider: authProviderLocalCassandra,sslOptions: { rejectUnauthorized: false },localDataCenter: 'West US'});</pre>
OPTIONAL: In package.json replace: "cassandra-driver": "datastax/nodejs-driver" with: "cassandra-driver": "^3.3.0"
WITH SSL-VALIDATION
<pre>var ssl_option = { cert : fs.readFileSync("/home/Cassandra-NodeJS-Demo/gabus2020cas.pem"), secureProtocol: 'TLSv1_2_method' }; var config = require('../config'); const authProviderLocalCassandra = new cassandra.auth.PlainTextAuthProvider(config.username, config.password); const client = new cassandra.Client({contactPoints: [config.contactPoint], authProvider: authProviderLocalCassandra, sslOptions: ssl_option,localDataCenter: 'West US'});</pre>
Remove: <pre>,function(nextCall) { client.execute("TRUNCATE nodejs_demo.us_unemployment", [],nextCall)}</pre>

Terminal-1 Node Environment
node src/Ingestion.js
node src/Reading.js
rm -r -f Cassandra-NodeJS-Demo/
rm -r -f azure-cosmos-db-cassandra-nodejs-getting-started/